

Chawin Sitawarin

✉ chawin.sitawarin@gmail.com
🌐 <https://chawin.sitawarin.com/>

Machine Learning Security & Privacy Researcher

Experiences

- 2026–present **Member of Technical Staff**, *Anthropic*, Safeguards Research.
- 2025–2026 **Research Scientist**, *Google DeepMind*, GenAI, Security and Privacy Research.
- 2024–2025 **Postdoctoral Researcher**, *Meta*, Central Applied Science, Privacy-Preserving Machine Learning. Memorization, privacy, and copyright risks in generative AI
- 2018–2024 **PhD in Computer Science**, *UC Berkeley*.
LLM security and adversarial robustness, advised by Prof. David Wagner
- 2014–2018 **BSE in Electrical Engineering (High Honor)**, *Princeton University*.
Certificate in Applications of Computing

Research Interests

I am broadly interested in security and privacy aspects of machine learning. My recent works are on **jailbreak** attacks and **prompt injection** defenses on large language models; my older works are on the **adversarial robustness** of machine learning algorithms. Currently, the problems I am excited about are:

- (1) Better evaluation of memorization, privacy, and copyright risks
- (2) Safety, security, and prompt injection defenses

Selected Publications

- 2026 **How Much Can Language Models Memorize?**, J. Morris, C. Sitawarin, C. Guo, N. Kokhlikyan, E. Suh, A. M. Rush, K. Chaudhuri, S. Mahloujifar, ICML 2026 (Oral, **Outstanding Paper Honorable Mention**) [paper](#).
- 2026 **Rapid Poison: Practical Poisoning Attacks Against the Rapid Response Framework**, D. Huang, J. Chang, A. Shah, P. Mittal, C. Sitawarin, ICML 2026 (Spotlight) [paper](#) [code](#).
- 2026 **The Attacker Moves Second: Stronger Adaptive Attacks Bypass Defenses Against LLM Jailbreaks and Prompt Injections**, M. Nasr*, N. Carlini*, C. Sitawarin*, S. V. Schulhoff*, J. Hayes, M. Ilie, J. Pluto, S. Song, H. Chaudhari, I. Shumailov, A. Thakurta, K. Y. Xiao, A. Terzis, F. Tramèr, USENIX Security 2026 [paper](#).
- 2025 **Stronger Universal and Transferable Attacks by Suppressing Refusals**, D. Huang, A. Shah, A. Araujo, D. Wagner, C. Sitawarin, NAACL 2025 [paper](#).
- 2025 **Mark My Words: Analyzing and Evaluating Language Model Watermarks**, J. Piet, C. Sitawarin, V. Fang, N. Mu, D. Wagner, SaTML 2025 [paper](#) [code](#).
- 2025 **Vulnerability Detection with Code Language Models: How Far Are We?**, Y. Ding, Y. Fu, O. Ibrahim, C. Sitawarin, X. Chen, B. Alomair, D. Wagner, B. Ray, Y. Chen, ICSE 2025 [paper](#) [code](#).
- 2025 **StruQ: Defending Against Prompt Injection with Structured Queries**, S. Chen, J. Piet, C. Sitawarin, D. Wagner, USENIX Security 2025 [paper](#) [code](#).
- 2024 **PAL: Proxy-Guided Black-Box Attack on Large Language Models**, C. Sitawarin, N. Mu, D. Wagner, A. Araujo, Preprint [paper](#) [code](#).
- 2024 **Jatmo: Prompt Injection Defense by Task-Specific Finetuning**, J. Piet*, M. Alrashed*, C. Sitawarin, et al., ESORICS 2024 [paper](#) [code](#).

- 2024 **PubDef: Defending against Transfer Attacks from Public Models**, C. Sitawarin, J. Chang*, D. Huang*, W. Altoyan, D. Wagner, ICLR 2024 (Poster) [paper](#) [code](#).
- 2023 **Preprocessors Matter! Realistic Decision-Based Attacks on Machine Learning Systems**, C. Sitawarin, F. Tramèr, N. Carlini, ICML 2023 (Poster) [paper](#) [code](#).
- 2023 **Part-Based Models Improve Adversarial Robustness**, C. Sitawarin, K. Pongmala, Y. Chen, N. Carlini, D. Wagner, ICLR 2023 (Poster) [paper](#) [code](#).
- 2022 **Demystifying the Adversarial Robustness of Random Transformation Defenses**, C. Sitawarin, Z. Golan-Strieb, D. Wagner, ICML 2022 and AAAI-22 AdvML Workshop (Best Paper) [paper](#) [code](#).
- 2021 **SAT: Improving Adversarial Training via Curriculum-Based Loss Smoothing**, C. Sitawarin, S. Chakraborty, D. Wagner, AISC 2021 (co-located with CCS) [paper](#).
- 2020 **Minimum-Norm Adversarial Examples on k -NN and k -NN-Based Models**, C. Sitawarin, D. Wagner, Deep Learning and Security Workshop (IEEE S&P 2020) [paper](#).
- 2018 **On the Robustness of Deep k -Nearest Neighbors**, C. Sitawarin, D. Wagner, Deep Learning and Security Workshop (IEEE S&P 2019) [paper](#).

Prior Internships

- Summer 2022 **Google**, Sunnyvale CA, Research intern.
Developed a defense against transfer adversarial attacks for a malware classification task with a pair of public client-side and secret server-side models. Hosted by Ali Zand and David Tao.
- Fall 2021–Spring 2022 **Google**, Remote, Part-time student researcher.
Developed new query-based adversarial attack and model-stealing attack against a black-box image preprocessing and recognition pipeline. Hosted by Nicholas Carlini.
- Summer 2021 **Nokia Bell Labs**, Remote, Research intern.
Investigated relationships between causality and robustness in machine learning, focusing on leveraging causal relationship to improve robustness and generalization to unseen corruptions. Hosted by Anwar Walid.
- Summer 2019 **IBM Research**, Yorktown Heights NY, Research intern.
Studied the effectiveness of existing defenses against adversarial examples from a perspective of concentration bound and improved adversarial training through optimization techniques. Hosted by Supriyo Chakraborty.

Awards & Grants

- 2023 **Outstanding Graduate Student Instructor Award** *Teaching award*
- 2022 **Google-BAIR Commons Project** *Research grant*
- 2021–2022 **Center for Long-Term Cybersecurity (CLTC)** *Research grant*
- 2021 **Microsoft-BAIR Commons Project** *Research grant*
- 2018 **Phi Beta Kappa** *Academic honor society*
- 2018 **Sigma Xi** *Scientific research honor society*
- 2017 **The P. Michael Lion III Fund** *Summer research funding for Princeton engineering students*
- 2016 **Tau Beta Pi** *Engineering honor society*
- 2016 **Shapiro Prize for Academic Excellence** *Academic award at Princeton University*
- 2013 **King's Scholarship** *Prestigious scholarship awarded by Thai government for pursuing a bachelor's degree*

Services

Reviewer, ICLR '24, '25 | ICML '22 (top reviewer), '24, '25, '26 | NeurIPS '22, '23, '24. '25 | TPAMI '24 | IEEE S&P '26 | CCS '26 | AISC '22, '23, '24, '25.